

1. Loading and Preprocessing

Load the breast cancer dataset from sklearn.

```
In [1]: from sklearn.datasets import load_breast_cancer
data = load_breast_cancer()
x = data.data
y = data.target
x,y
```

```
Out[1]: (array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
1.189e-01],
[2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
8.902e-02],
[1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
8.758e-02],
...,
[1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
7.820e-02],
[2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
1.240e-01],
[7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
7.039e-02]]),
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0,
1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0,
1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1,
0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0,
0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0,
1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1,
1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0,
0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,
0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0,
1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1]))
```

Preprocess the data to handle any missing values and perform necessary feature scaling.

```
In [2]: import numpy as np
print(np.isnan(x).sum())
```

0

Explain the preprocessing steps you performed and justify why they are necessary for this dataset.

```
In [4]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
```

```
In [5]: x_scaled
```

```
Out[5]: array([[ 1.09706398, -2.07333501,  1.26993369, ...,  2.29607613,
                2.75062224,  1.93701461],
               [ 1.82982061, -0.35363241,  1.68595471, ...,  1.0870843 ,
               -0.24388967,  0.28118999],
               [ 1.57988811,  0.45618695,  1.56650313, ...,  1.95500035,
                1.152255  ,  0.20139121],
               ...,
               [ 0.70228425,  2.0455738 ,  0.67267578, ...,  0.41406869,
               -1.10454895, -0.31840916],
               [ 1.83834103,  2.33645719,  1.98252415, ...,  2.28998549,
                1.91908301,  2.21963528],
               [-1.80840125,  1.22179204, -1.81438851, ..., -1.74506282,
               -0.04813821, -0.75120669]])
```

2. Classification Algorithm Implementation



1. Logistic Regression

```
In [7]: from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression(max_iter=10000)
log_reg.fit(x_scaled, y)
```

```
Out[7]:
```

▼	LogisticRegression
	LogisticRegression(max_iter=10000)

2. Decision Tree Classifier

```
In [9]: from sklearn.tree import DecisionTreeClassifier  
tree_clf = DecisionTreeClassifier()  
tree_clf.fit(x_scaled, y)
```

```
Out[9]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

3. Random Forest Classifier

```
In [12]: from sklearn.ensemble import RandomForestClassifier  
rf_clf = RandomForestClassifier()  
rf_clf.fit(x_scaled, y)
```

```
Out[12]: ▾ RandomForestClassifier  
RandomForestClassifier()
```

4. Support Vector Machine (SVM)

```
In [13]: from sklearn.svm import SVC  
svm_clf = SVC()  
svm_clf.fit(x_scaled, y)
```

```
Out[13]: ▾ SVC  
SVC()
```

5. k-Nearest Neighbors (k-NN)

```
In [14]: from sklearn.neighbors import KNeighborsClassifier  
knn_clf = KNeighborsClassifier()  
knn_clf.fit(x_scaled, y)
```

```
Out[14]: ▾ KNeighborsClassifier  
KNeighborsClassifier()
```

3. Model Comparison

```
In [15]: from sklearn.model_selection import cross_val_score
models = [log_reg, tree_clf, rf_clf, svm_clf, knn_clf]
model_names = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM']
for model, name in zip(models, model_names):
    scores = cross_val_score(model, x_scaled, y, cv=5, scoring='accuracy')
    print(f"{name}: {scores.mean():.4f}")
```

Logistic Regression: 0.9807

Decision Tree: 0.9138

Random Forest: 0.9631

SVM: 0.9736

k-NN: 0.9649

Which algorithm performed the best and which one performed the worst?

Best Algorithm: The one with the highest average accuracy. Worst Algorithm: The one with the lowest average accuracy.

In []: